

Understanding the Data

```
In [ ]: import os
import tarfile
from six.moves import urllib
import pandas as pd

DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml2/master/"
HOUSING_PATH = os.path.join("datasets", "housing")
HOUSING_URL = DOWNLOAD_ROOT + "datasets/housing/housing.tgz"

def fetch_housing_data(housing_url=HOUSING_URL, housing_path=HOUSING_PATH):
    if not os.path.isdir(housing_path):
        os.makedirs(housing_path)
    tgz_path = os.path.join(housing_path, "housing.tgz")
    urllib.request.urlretrieve(housing_url, tgz_path)
    housing_tgz = tarfile.open(tgz_path)
    housing_tgz.extractall(path=housing_path)
    housing_tgz.close()
```

upon calling the function `fetch_housing_data()` , it

1. creates datasets/housing directory
2. downloads housing.tgz
3. extracts housing.csv in the directory

```
In [ ]: def load_housing_data(housing_path=HOUSING_PATH):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)
```

```
In [ ]: fetch_housing_data() # get the csv

housing = load_housing_data() # load it
housing.head()
```

```
Out [ ]: 
```

| | longitude | latitude | housing_median_age | total_rooms | total_bedrooms | population | households | median_incor |
|---|-----------|----------|--------------------|-------------|----------------|------------|------------|--------------|
| 0 | -122.23 | 37.88 | 41.0 | 880.0 | 129.0 | 322.0 | 126.0 | 8.32 |
| 1 | -122.22 | 37.86 | 21.0 | 7099.0 | 1106.0 | 2401.0 | 1138.0 | 8.30 |
| 2 | -122.24 | 37.85 | 52.0 | 1467.0 | 190.0 | 496.0 | 177.0 | 7.25 |
| 3 | -122.25 | 37.85 | 52.0 | 1274.0 | 235.0 | 558.0 | 219.0 | 5.64 |
| 4 | -122.25 | 37.85 | 52.0 | 1627.0 | 280.0 | 565.0 | 259.0 | 3.84 |

```
In [ ]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
In [ ]: # examine the ocean_proximity column
```

```
housing["ocean_proximity"].value_counts()
```

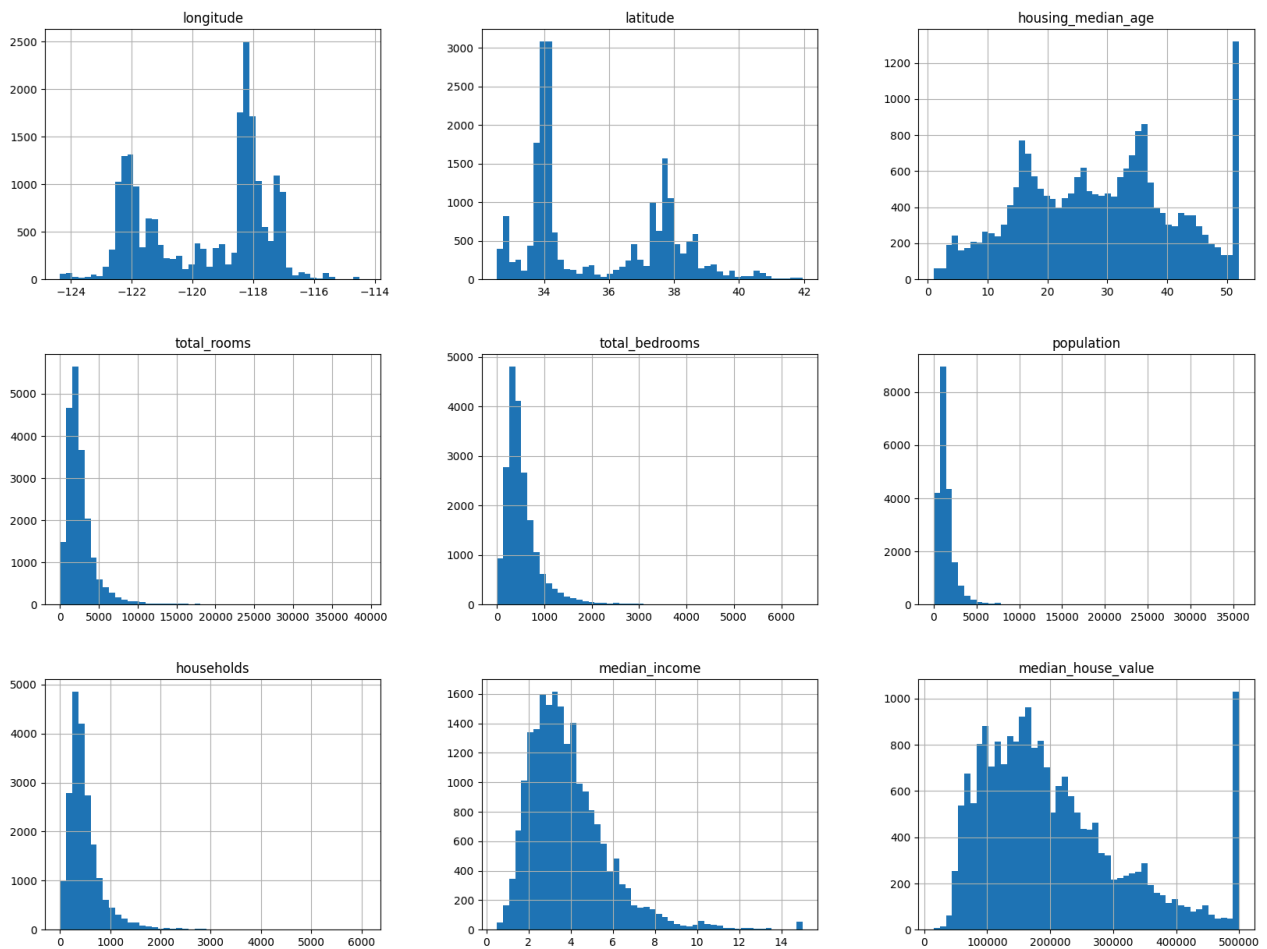
```
Out[ ]: ocean_proximity
<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
Name: count, dtype: int64
```

```
In [ ]: housing.describe() # describes the summary of numerical attributes
```

```
Out[ ]:
```

| | longitude | latitude | housing_median_age | total_rooms | total_bedrooms | population | households |
|--------------|--------------|--------------|--------------------|--------------|----------------|--------------|--------------|
| count | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20433.000000 | 20640.000000 | 20640.000000 |
| mean | -119.569704 | 35.631861 | 28.639486 | 2635.763081 | 537.870553 | 1425.476744 | 499.539000 |
| std | 2.003532 | 2.135952 | 12.585558 | 2181.615252 | 421.385070 | 1132.462122 | 382.329000 |
| min | -124.350000 | 32.540000 | 1.000000 | 2.000000 | 1.000000 | 3.000000 | 1.000000 |
| 25% | -121.800000 | 33.930000 | 18.000000 | 1447.750000 | 296.000000 | 787.000000 | 280.000000 |
| 50% | -118.490000 | 34.260000 | 29.000000 | 2127.000000 | 435.000000 | 1166.000000 | 409.000000 |
| 75% | -118.010000 | 37.710000 | 37.000000 | 3148.000000 | 647.000000 | 1725.000000 | 605.000000 |
| max | -114.310000 | 41.950000 | 52.000000 | 39320.000000 | 6445.000000 | 35682.000000 | 6082.000000 |

```
In [ ]: housing.hist(bins=50, figsize=(20,15))
__import__("matplotlib").pyplot.show()
```



Vertical Axis → Number of entries/instances of the value

Horizontal Axis → range of values

```
In [ ]: housing["median_income"]
```

```
Out [ ]: 0      8.3252
         1      8.3014
         2      7.2574
         3      5.6431
         4      3.8462
         ...
        20635   1.5603
        20636   2.5568
        20637   1.7000
        20638   1.8672
        20639   2.3886
        Name: median_income, Length: 20640, dtype: float64
```

You find out that the median_income is represented in ten thousands of US dollars.
eg: 8 means \$80,000

Making a Test Set and Training Set

```
In [ ]: # creating a test set (generally 20% of the dataset or even less if dataset is too large)

import numpy as np

def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data) * test_ratio)
    test_indices, train_indices = shuffled_indices[:test_set_size], shuffled_indices[test_set_size:]
```

```

    return data.iloc[test_indices], data.iloc[train_indices]

test_set, train_set = split_train_test(housing, 0.2)
test_set.head()

```

Out []:

| | longitude | latitude | housing_median_age | total_rooms | total_bedrooms | population | households | median_i |
|-------|-----------|----------|--------------------|-------------|----------------|------------|------------|----------|
| 1121 | -121.58 | 39.79 | 19.0 | 2636.0 | 523.0 | 1184.0 | 465.0 | |
| 14206 | -117.06 | 32.69 | 9.0 | 521.0 | 111.0 | 491.0 | 110.0 | |
| 5256 | -118.48 | 34.07 | 37.0 | 4042.0 | 549.0 | 1318.0 | 542.0 | |
| 15149 | -116.90 | 32.90 | 19.0 | 3090.0 | 552.0 | 1621.0 | 520.0 | |
| 16915 | -122.35 | 37.56 | 52.0 | 1659.0 | 191.0 | 519.0 | 201.0 | |

Break-proofing the split approach

The next time the code is run again, it will choose different set of indices, over time, the model will get to see the whole dataset which is to be avoided.

Solutions:

1. Save test set on first run and load it in subsequent runs
2. Set a random number generator's seed before calling permutation() so it generates same shuffled indices

But both the solutions will break when fetching an updated dataset.

```

In [ ]: # Best solution

from zlib import crc32

def test_set_check(identifier, test_ratio):
    return crc32(np.int64(identifier)) & 0xffffffff < test_ratio * 2 ** 32

def split_train_test_by_id(data, test_ratio, id_column_name):
    ids = data[id_column_name]
    in_test_set = ids.apply(lambda id_: test_set_check(id_, test_ratio))
    return data.loc[~in_test_set], data.loc[in_test_set]

```

Working of the test_set_check

1. the identifier is converted into a 64-bit integer.
2. The crc32 function returns the hash of the 64-bit identifier and it is masked with the highest value of 32 bit integer which is 0xffffffff to truncate it into a 32-bit integer.
3. the test_ratio is multiplied with 2^{32} to maintain the scale.
4. The hash is checked with the test_ratio

```

In [ ]: # Add a custom ID column to the dataset since it does not have it

housing

```

Out[]:

| | longitude | latitude | housing_median_age | total_rooms | total_bedrooms | population | households | median_i |
|-------|-----------|----------|--------------------|-------------|----------------|------------|------------|----------|
| 0 | -122.23 | 37.88 | 41.0 | 880.0 | 129.0 | 322.0 | 126.0 | |
| 1 | -122.22 | 37.86 | 21.0 | 7099.0 | 1106.0 | 2401.0 | 1138.0 | |
| 2 | -122.24 | 37.85 | 52.0 | 1467.0 | 190.0 | 496.0 | 177.0 | |
| 3 | -122.25 | 37.85 | 52.0 | 1274.0 | 235.0 | 558.0 | 219.0 | |
| 4 | -122.25 | 37.85 | 52.0 | 1627.0 | 280.0 | 565.0 | 259.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 20635 | -121.09 | 39.48 | 25.0 | 1665.0 | 374.0 | 845.0 | 330.0 | |
| 20636 | -121.21 | 39.49 | 18.0 | 697.0 | 150.0 | 356.0 | 114.0 | |
| 20637 | -121.22 | 39.43 | 17.0 | 2254.0 | 485.0 | 1007.0 | 433.0 | |
| 20638 | -121.32 | 39.43 | 18.0 | 1860.0 | 409.0 | 741.0 | 349.0 | |
| 20639 | -121.24 | 39.37 | 16.0 | 2785.0 | 616.0 | 1387.0 | 530.0 | |

20640 rows × 10 columns



In []: `housing_with_id = housing.reset_index() # adds an index column`
`housing_with_id`

Out[]:

| | index | longitude | latitude | housing_median_age | total_rooms | total_bedrooms | population | households | n |
|-------|-------|-----------|----------|--------------------|-------------|----------------|------------|------------|-----|
| 0 | 0 | -122.23 | 37.88 | 41.0 | 880.0 | 129.0 | 322.0 | 126.0 | |
| 1 | 1 | -122.22 | 37.86 | 21.0 | 7099.0 | 1106.0 | 2401.0 | 1138.0 | |
| 2 | 2 | -122.24 | 37.85 | 52.0 | 1467.0 | 190.0 | 496.0 | 177.0 | |
| 3 | 3 | -122.25 | 37.85 | 52.0 | 1274.0 | 235.0 | 558.0 | 219.0 | |
| 4 | 4 | -122.25 | 37.85 | 52.0 | 1627.0 | 280.0 | 565.0 | 259.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 20635 | 20635 | -121.09 | 39.48 | 25.0 | 1665.0 | 374.0 | 845.0 | 330.0 | |
| 20636 | 20636 | -121.21 | 39.49 | 18.0 | 697.0 | 150.0 | 356.0 | 114.0 | |
| 20637 | 20637 | -121.22 | 39.43 | 17.0 | 2254.0 | 485.0 | 1007.0 | 433.0 | |
| 20638 | 20638 | -121.32 | 39.43 | 18.0 | 1860.0 | 409.0 | 741.0 | 349.0 | |
| 20639 | 20639 | -121.24 | 39.37 | 16.0 | 2785.0 | 616.0 | 1387.0 | 530.0 | |

20640 rows × 11 columns



In []: `train_set, test_set = split_train_test_by_id(housing_with_id, 0.2, "index")`
`test_set`

Out[]:

| | index | longitude | latitude | housing | median_age | total_rooms | total_bedrooms | population | households | n |
|--|-------|-----------|----------|---------|------------|-------------|----------------|------------|------------|---|
| | 2 | -122.24 | 37.85 | | 52.0 | 1467.0 | 190.0 | 496.0 | 177.0 | |
| | 5 | -122.25 | 37.85 | | 52.0 | 919.0 | 213.0 | 413.0 | 193.0 | |
| | 12 | -122.26 | 37.85 | | 52.0 | 2491.0 | 474.0 | 1098.0 | 468.0 | |
| | 16 | -122.27 | 37.85 | | 52.0 | 1966.0 | 347.0 | 793.0 | 331.0 | |
| | 23 | -122.27 | 37.84 | | 52.0 | 1688.0 | 337.0 | 853.0 | 325.0 | |
| | ... | ... | ... | | ... | ... | ... | ... | ... | |
| | 20615 | -121.54 | 39.08 | | 23.0 | 1076.0 | 216.0 | 724.0 | 197.0 | |
| | 20617 | -121.53 | 39.06 | | 20.0 | 561.0 | 109.0 | 308.0 | 114.0 | |
| | 20622 | -121.44 | 39.00 | | 20.0 | 755.0 | 147.0 | 457.0 | 157.0 | |
| | 20626 | -121.43 | 39.18 | | 36.0 | 1124.0 | 184.0 | 504.0 | 171.0 | |
| | 20629 | -121.39 | 39.12 | | 28.0 | 10035.0 | 1856.0 | 6912.0 | 1818.0 | |

4128 rows × 11 columns



Same thing can be done with Scikit-Learn's train_test_split method:

In []:

```
from sklearn.model_selection import train_test_split

train_test, test_set = train_test_split(housing, test_size=0.2, random_state=42) # the integer 42 is
test_set # run it as many times as you want and the result will be the same (unless you change the
```

Out[]:

| | longitude | latitude | housing | median_age | total_rooms | total_bedrooms | population | households | median_i |
|-------|-----------|----------|---------|------------|-------------|----------------|------------|------------|----------|
| 20046 | -119.01 | 36.06 | | 25.0 | 1505.0 | NaN | 1392.0 | 359.0 | |
| 3024 | -119.46 | 35.14 | | 30.0 | 2943.0 | NaN | 1565.0 | 584.0 | |
| 15663 | -122.44 | 37.80 | | 52.0 | 3830.0 | NaN | 1310.0 | 963.0 | |
| 20484 | -118.72 | 34.28 | | 17.0 | 3051.0 | NaN | 1705.0 | 495.0 | |
| 9814 | -121.93 | 36.62 | | 34.0 | 2351.0 | NaN | 1063.0 | 428.0 | |
| ... | ... | ... | | ... | ... | ... | ... | ... | |
| 15362 | -117.22 | 33.36 | | 16.0 | 3165.0 | 482.0 | 1351.0 | 452.0 | |
| 16623 | -120.83 | 35.36 | | 28.0 | 4323.0 | 886.0 | 1650.0 | 705.0 | |
| 18086 | -122.05 | 37.31 | | 25.0 | 4111.0 | 538.0 | 1585.0 | 568.0 | |
| 2144 | -119.76 | 36.77 | | 36.0 | 2507.0 | 466.0 | 1227.0 | 474.0 | |
| 3665 | -118.37 | 34.22 | | 17.0 | 1787.0 | 463.0 | 1671.0 | 448.0 | |

4128 rows × 10 columns




Avoiding Sampling Bias

It is important for a model to consider each and every category of an attribute to avoid sampling bias. for example, a survey on population of a country consisting of 55% males and 45% females need to take into consideration of 55 males and 45 females in a 100 people. If the percentages does not represent the whole population of the country, then the survey may be considered as biased.

Suppose the experts told you that the median income is a very important attribute to predict median housing prices. You should ensure that the test set is representative of the various categories of incomes in the whole dataset.

```
In [ ]: plt = __import__("matplotlib").pyplot # using matplotlib's pyplot
```

```
In [ ]: housing["median_income"].hist(figsize=(5, 5), bins=50)
plt.show()
```



```
In [ ]: housing["income_cat"] = pd.cut(
    housing["median_income"],
    bins=[0, 1.5, 3.0, 4.5, 6, np.inf],
    labels=[1, 2, 3, 4, 5]
)

housing[["median_income", "income_cat"]]
```

```
Out[ ]:
```


| | median_income | income_cat |
|--|---------------|------------|
|--|---------------|------------|

| | | |
|-------|--------|-----|
| 0 | 8.3252 | 5 |
| 1 | 8.3014 | 5 |
| 2 | 7.2574 | 5 |
| 3 | 5.6431 | 4 |
| 4 | 3.8462 | 3 |
| ... | ... | ... |
| 20635 | 1.5603 | 2 |
| 20636 | 2.5568 | 2 |
| 20637 | 1.7000 | 2 |
| 20638 | 1.8672 | 2 |
| 20639 | 2.3886 | 2 |

20640 rows × 2 columns

```
In [ ]: housing["income_cat"].hist()
```

```
Out [ ]: <Axes: >
```



Stratified Sampling

```
In [ ]: from sklearn.model_selection import StratifiedShuffleSplit

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_indexes, test_indexes in split.split(housing, housing["income_cat"]):
    strat_train_set = housing.loc[train_indexes]
    strat_test_set = housing.loc[test_indexes]

strat_train_set
```

```
Out [ ]:
```

| | longitude | latitude | housing_median_age | total_rooms | total_bedrooms | population | households | median_i |
|--------------|-----------|----------|--------------------|-------------|----------------|------------|------------|----------|
| 12655 | -121.46 | 38.52 | 29.0 | 3873.0 | 797.0 | 2237.0 | 706.0 | |
| 15502 | -117.23 | 33.09 | 7.0 | 5320.0 | 855.0 | 2015.0 | 768.0 | |
| 2908 | -119.04 | 35.37 | 44.0 | 1618.0 | 310.0 | 667.0 | 300.0 | |
| 14053 | -117.13 | 32.75 | 24.0 | 1877.0 | 519.0 | 898.0 | 483.0 | |
| 20496 | -118.70 | 34.28 | 27.0 | 3536.0 | 646.0 | 1837.0 | 580.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 15174 | -117.07 | 33.03 | 14.0 | 6665.0 | 1231.0 | 2026.0 | 1001.0 | |
| 12661 | -121.42 | 38.51 | 15.0 | 7901.0 | 1422.0 | 4769.0 | 1418.0 | |
| 19263 | -122.72 | 38.44 | 48.0 | 707.0 | 166.0 | 458.0 | 172.0 | |
| 19140 | -122.70 | 38.31 | 14.0 | 3155.0 | 580.0 | 1208.0 | 501.0 | |
| 19773 | -122.14 | 39.97 | 27.0 | 1079.0 | 222.0 | 625.0 | 197.0 | |

16512 rows × 11 columns

```
In [ ]: # percentage of categories in train set
highest = max([j for _, j in strat_train_set["income_cat"].value_counts().items()])
```



```
for i,j in strat_train_set["income_cat"].value_counts().items():
    print(str(i) + ":", str((j/highest)*100) + "%")
```

```
3: 100.0%
2: 90.9483503195716%
4: 50.28502332008983%
5: 32.64812575574365%
1: 11.349110381758507%
```

```
In [ ]: # percentage of categories in test set
highest = max([j for _,j in strat_test_set["income_cat"].value_counts().items()])
for i,j in strat_test_set["income_cat"].value_counts().items():
    print(str(i) + ":", str((j/highest)*100) + "%")
```

```
3: 100.0%
2: 90.94678645473392%
4: 50.310988251554946%
5: 32.61921216309606%
1: 11.402902557014514%
```

Since both the train set and test set have the same number of categories shuffled randomly, the likelihood of a Sampling Bias is significantly reduced.

You should remove the income_cat attribute so the data is back to its original state.

```
In [ ]: # Removing income_cat column
for set_ in (strat_train_set, strat_test_set):
    set_.drop("income_cat", axis=1, inplace=True)
```

Visualize the Data

The goal is to go a little more in-depth into understanding of the data. First, make sure the test-set is left aside and only the training set is explored.

Note: If the training set is very large, you may want to sample an exploration set, to make manipulations easy and fast. In this case, the training set is fairly small so you can directly work on the full set.

```
In [ ]: # Plotting latitudes and longitudes to see the location of all districts

housing.plot(kind="scatter", x="longitude", y="latitude")
plt.show()
```





fig: scattered districts of California


```
In [ ]: housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
plt.show()
```



highly densed areas can now be seen clearly.

Now let's visualize the population of each district in the state. The radius of the circles represents the district's population and the color represents the price (blue meaning more affordable and red being more expensive).

```
In [ ]: housing.plot(
    kind="scatter",
    alpha=0.3,
    x="longitude",
    y="latitude",
    s=housing["population"]/100,
    c="median_house_value",
    cmap=plt.get_cmap("jet"),
    colorbar=True,
    figsize=(10,7)
)
plt.show()
```



Takeaways from the graph:

1. The housing prices are related to the location(eg, close to the ocean)
2. The housing prices are related to population density

Looking for Correlations

We can compute the standard correlation coefficient between every pair of attributes using the `corr()` method.

```
In [ ]: # housing_numeric_only = housing.select_dtypes(include=[np.number])
# corr_matrix = housing_numeric_only.corr()
# corr_matrix["median_house_value"].sort_values(ascending=False)

# (or)

corr_matrix = housing.corr(numeric_only=True)
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
Out [ ]: median_house_value    1.000000
median_income      0.688075
total_rooms        0.134153
housing_median_age  0.105623
households         0.065843
total_bedrooms     0.049686
population         -0.024650
longitude          -0.045967
latitude           -0.144160
Name: median_house_value, dtype: float64
```

These correlations co-efficients are with respect the the `median_house_value` . For example, the `median_house_value` tends to go up when the `median_income` goes up, because the correlation coefficient between the two is 0.688, which is pretty close to 1 and has a strong positive correlation.

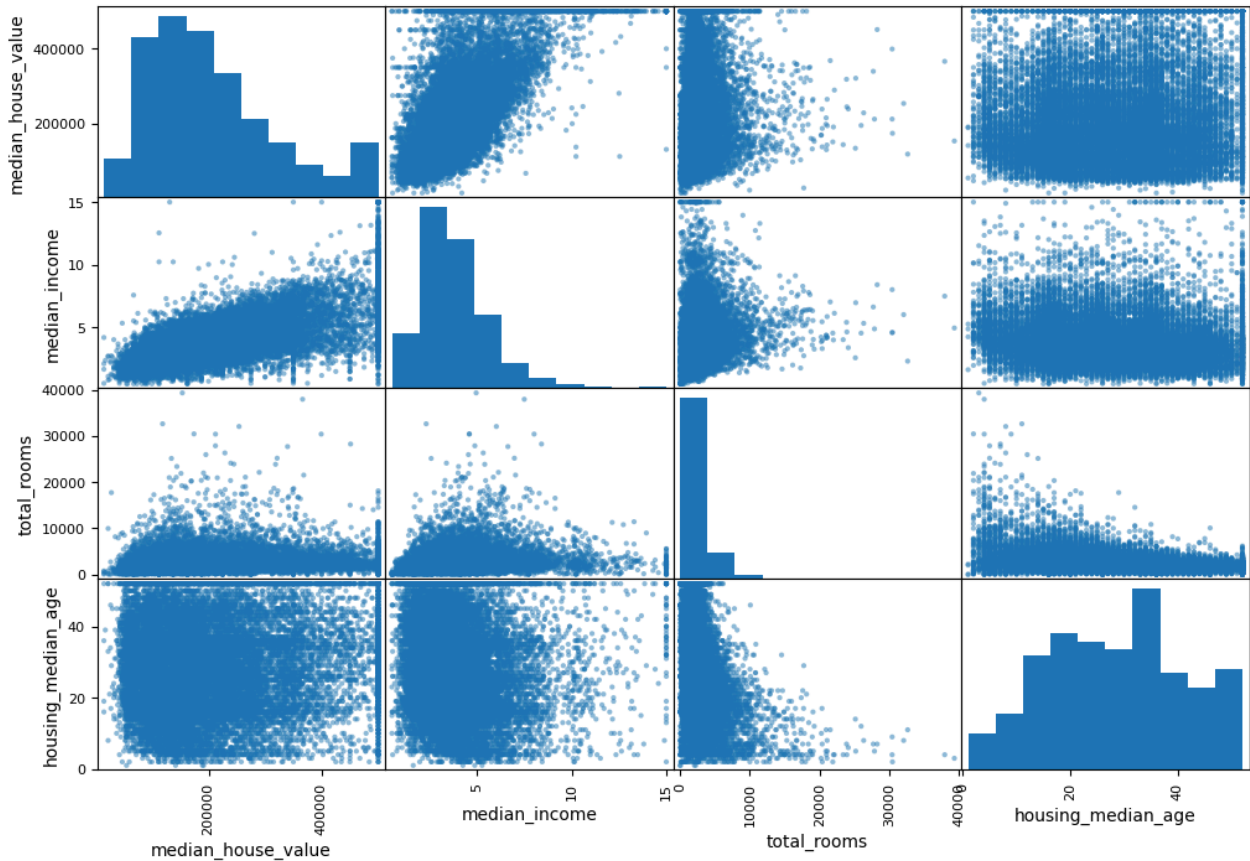
Also, the `median_house_value` tends to go a little up the more you travel to the south because of the negative correlation between `median_house_value` and `latitude`.

Another way is to use `pandas.plotting.scatter_matrix` function, which plots every numerical attribute against every other numerical attribute. This results in $11^2 = 121$ correlations and 121 plots which won't fit in the graph. So we will only focus on a few promising attributes that seem most correlated with the `median_house_value`.

```
In [ ]: from pandas.plotting import scatter_matrix

attributes = ["median_house_value", "median_income", "total_rooms", "housing_median_age"]

_ = scatter_matrix(housing[attributes], figsize=(12, 8))
plt.show()
```



Since the most useful attribute is the `median_income` attribute, we will focus on that.

```
In [ ]: housing.plot(kind="scatter", x="median_income", y="median_house_value", alpha=0.1, figsize=(8,4))
plt.show()
```

